

# 3 Mini Talks of 10 minutes each



**OS/application inversion**  
Anil Madhavapeddy



**Living Databases**  
Pavlo Baron



**Why do we still not give a  
@#%! about testing?!**  
Trisha Gee

The OS/App Inversion:  
escaping POSIX to bring Git to the datacentre  
QCon New York Mini talk

Anil Madhavapeddy (speaker)  
with Thomas Gazagnaire and Benjamin Farinier  
University of Cambridge Computer Laboratory

June 7, 2015

# Common features every distributed system needs

- **Persistence** for fault tolerance and scaling
- **Scheduling** of communication between nodes
- **Tracing** across nodes for debugging and profiling

Most distributed systems run over an operating system, and so are stuck with the OS kernel exerting control. We use *unikernels*, which are application VMs that have complete control over their resources.

# What if we just used Git?

- **Persistence**

- `git clone` of a shared repository across nodes
- `git commit` of local operations in the node

- **Scheduling**

- `git pull` to receive events from other nodes
- `git push` to publish events to other nodes

- **Tracing and Debugging**

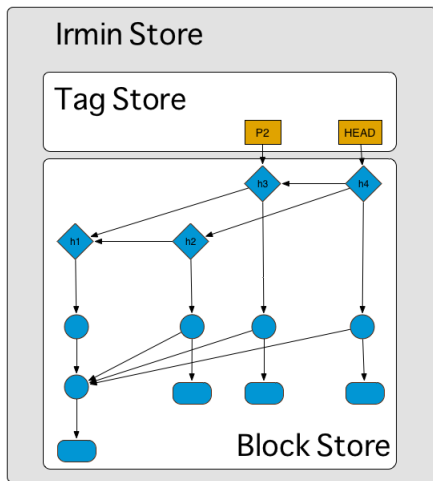
- `git log` to see global operations
- `git checkout` to roll back time to a snapshot
- `git bisect` to locate problem messages

- **New Problems**

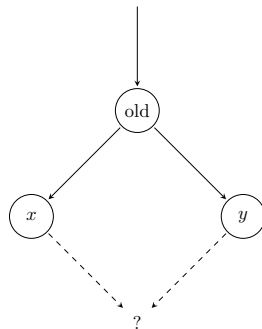
- `git rebase` needed for distributed garbage collection
- Shelling out to `git` is slow and lacks control

# Irmin, large-scale, immutable, branch-consistent storage

- Irmin is a library to **persist** and **synchronise distributed data structures** both on-disk and in-memory
- It enables a style of programming very similar to the **Git workflow**, where distributed **nodes fork, fetch, merge and push** data between each other
- The general idea is that you want every active node to get a **local (partial) copy of a global database** and always be very explicit about how and when data is shared and migrated



```
type t = ...  
(** User-defined contents. *)  
type result = [ `Ok of t |  
  `Conflict of string ]  
  
val merge: old:t → t → t →  
  result  
(** 3-way merge functions. *)
```



# Irmin Features

- Still pre 1.0, but several useful datastructures such as distributed queues and efficient ropes.
- HTTP REST for remote clients, or library via OCaml.
- JavaScript compilation for pure browser operation.
- Bidirectional operation, so `git` commits map to Irmin commits from any direction.
- Open source at <https://irmin.io>
- **Want to know more?** I'm giving a full talk on this on Friday at 1015 titled "Functional Distributed Programming with Irmin"!



## Teaser: Xen Toolstack using Irmin

<https://www.youtube.com/watch?v=DSzvFwIVm5s>

# Why Do We Still Not We Give A @#\$! About Testing?

Trisha Gee @trisha\_gee

Developer Advocate at JetBrains and OCD Geek

# 2010 IT Project Success Rates

Project Style	Successful	Challenged	Failure
Ad-hoc	49%	37%	14%
Iterative	61%	28%	11%
Agile	60%	28%	12%
Traditional	47%	36%	17%

[Dr Dobbs, 2010](#)

...when it comes to quality, 40% prefer to deliver on time and on budget and 57% prefer to deliver high-quality, easy-to-maintain systems.

# Software Project Success Rates

Successful	32%
Challenged	44%
Failure	24%

[Gartner, 2012](#)

So, Why Don't We Test?

Focus is on  
delivering features

Boring

Too much effort for  
too little value

"I don't know what  
it's supposed to do"

Code wasn't designed  
for testing

Too Hard

Just a box  
to tick

Framework  
gets in the way

"It's only a prototype"

Takes too  
much time

I'm the only  
maintainer

The code is  
self describing

Manual  
Process

Why Should We Care?



Documentation Freedom to  
Refactor

Knowledge  
Sharing

Usability

Help us  
with Design

So we know  
what doesn't  
work

Identify missing  
features

So we know  
it works!

What Can We Do?

# Tap into Motivation

- OCD
- Frame it as a difficult problem
- Gamification
- Make it as easy as possible
- Lead by example
- Exploration (what if...?)
- Set as primary task
- Add to Objectives
- Break emotional attachment to code

We Don't Have To Care About  
Testing

We Need To Encourage  
Behaviour That Prioritises Quality

# “living” database

@pavlobaron

speaking of  
streaming, reactive,  
realtime, near time,  
near realtime etc.

here is how all things are  
supposed to look in the  
reactive context







here is how current  
database access methods  
look in the reactive context







here is how current  
databases look in the  
reactive context





9322537031465





it's time to  
rethink databases

“living” database

some of the ideas are  
partially implemented  
in CEP, NoSQL, classic  
DBMS, Rx etc.,  
but not as a whole



storage is an aspect

storage is raw

everything is  
ordered by time

everything is on channels

queries are continuous

views are materialised  
and continuous

results are published

results are limited,  
timed-out or ignored  
by consumer



database quantifies harvest

changes are published  
both ways

everybody fires and forgets,  
but gets reminded  
when things go wrong

consumers react

it needs far more than  
just support in  
the database driver

it needs the database  
to be an active part  
of the overall data flow

“living” database

# Questions